



Calibration decides whether the detector works: the raw clean statistic is $8\times$ heavier-tailed than Gaussian at 3σ , the read-out decides which attacks are visible at all, and the nominal 5% false-positive rate reads 5 to 10% held out and is unusable at 25 prompts.



(b) GPT-2 layer 8: three read-outs of one inversion



(c) the nominal rate is a training constraint

